Crash Conjectures

Applied Project

Group 1: Unsupervised Twitter Robots

Dagny Stahl ◆ Luke Valentine ◆ Elise Renwick

Purpose

Create a model to predict if an accident, based on the current driving situation, would cause casualties in order to inform self driving cars

Product

API that a self driving car can interact with!

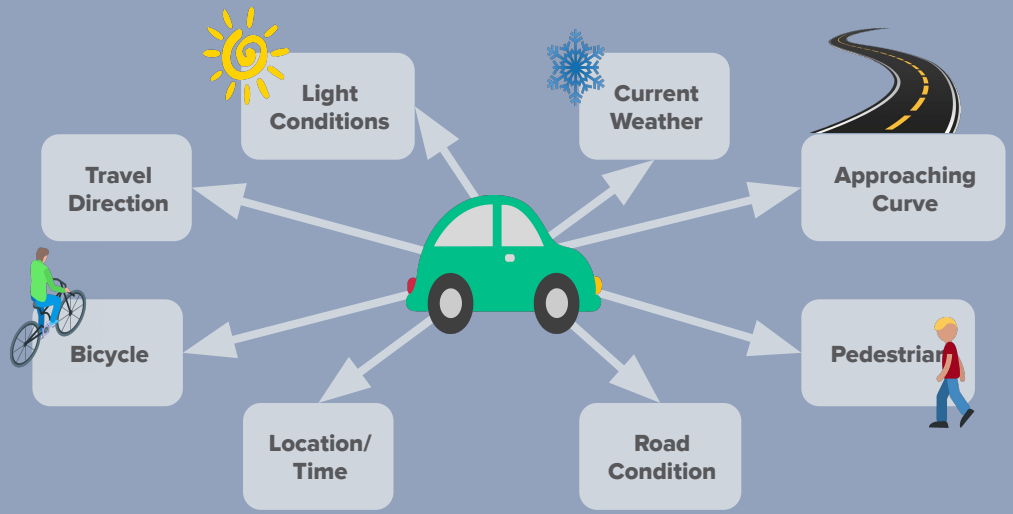

Sends Request

```
{
    "casualties": false,
}
```

Demo at end

The Main Dataset

Denver Car accidents data from CDOT, 189k rows

Information like:

- car type
- road & light conditions
- pedestrian or bike involvement
- Location
- date/time
- car movement

Mix of numerical, text, date, and time data

Feature Selection

Mutual Information Score Feature Selection

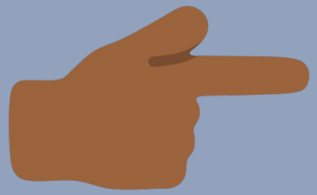

BICYCLE_IND is **0.077**
PEDESTRIAN_IND is **64.729**
ACCIDENT_OFFENSE_ACCIDENT is **12.36**
ACCIDENT_OFFENSE_DUI/DUID is **2.85**
NEIGHBORHOOD_Congress_Park is **6.53**
HIGHWAY_INTERCHANGE is **12.12**
INTERSECTION is **0.556**
PARKING_LOT is **0.466**
ROAD_CONTOUR_HILLCREST is **0.025**
ROAD_CONTOUR_STRAIGHT is **1.22**
ROAD_CONDITION_DRY is **1.33**
ROAD_CONDITION_ICY is **0.32**
LIGHT_CONDITION_DARKis **5.1**
LIGHT_CONDITION_DAY_LIGHT is **3.50**
DIRECTION_NORTH is **1.78**
DIRECTION_NORTHEAST is **0.051**

Data Cleaning

Data Formatting:

- Got date and time in a consistent format
- Dropped unnecessary columns (ex.“precinct ID”, “latitude”)
- Combined “fatalities” and “injuries” to “casualties”
- Ran dummies on categorical columns (ex. “neighborhood”, “road_condition”)
- Balanced data (1.7% of samples had casualties)
- Normalized the data
- Kept features a self driving car would have, (ex. removed “accident_offense_DUI”

Feature Selection

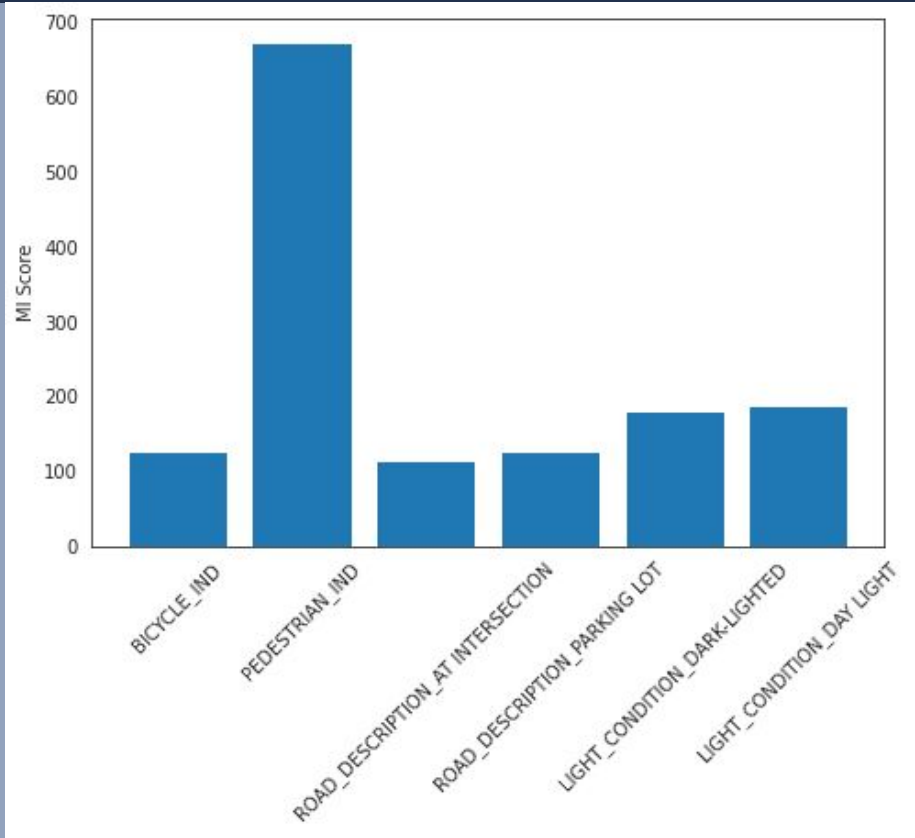

Feature Exploration

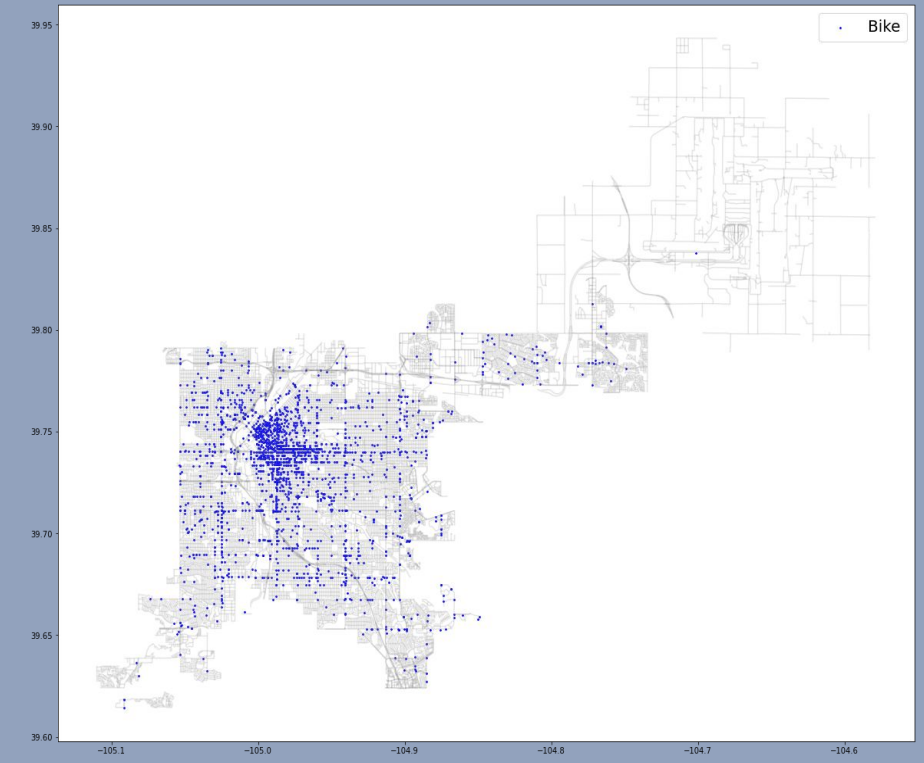

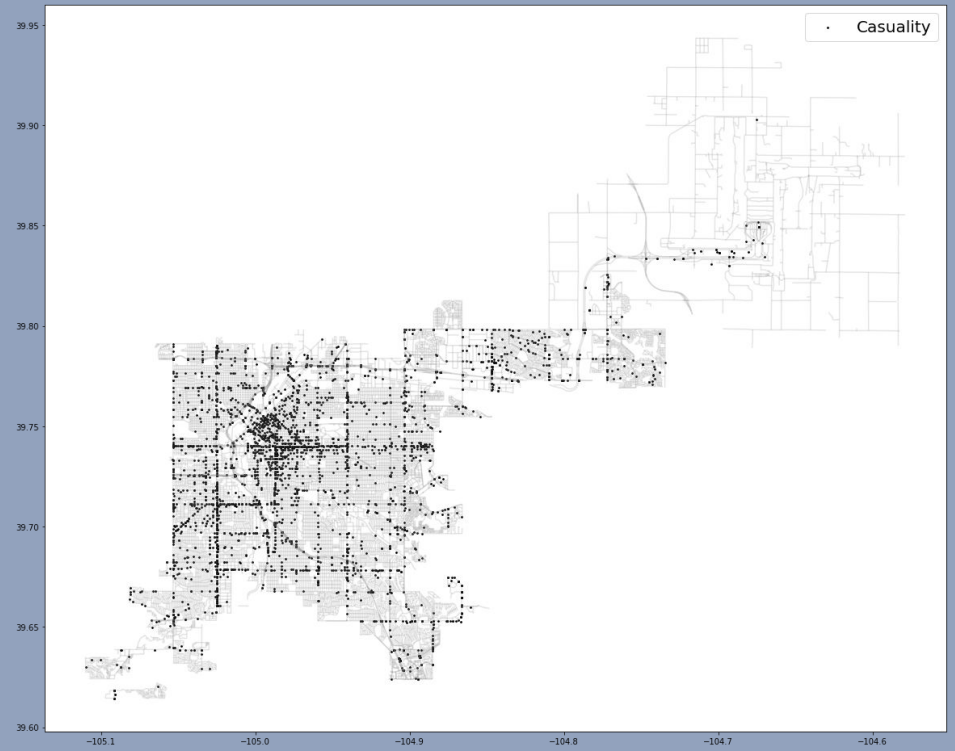

Failed Models: Regression

- Linear
- KNN
- SGD

Scored well on unbalanced data:
R^2 > 0.84 & MSE < 0.0106

...But horribly on balanced data:
R^2 < 0

The Problem: Only predicting no casualties, which was a safe bet, but the model was useless

Failed Models: Neural Net

New NN
Chose 20 best features from MI:

- NEIGHBORHOOD_ID
- BICYCLE_IND
- PEDESTRIAN_IND
- ROAD_CONDITION_ICY
- LIGHT_CONDITION_DUSK
- TIME_HR
- ...

to predict: CASUALTIES

Layers:

Dense(10,"relu")
Dropout(.9)
Dense(10,'relu')
Dropout(.9)
Dense(1)

Issues with our first NN

- **Trained on unbalanced data**
- **Predicting fatalities from info such as injuries**
- **Used features a SDC wouldn't have**

Failed Models: Neural Net

Scored well on unbalanced data with more features:
loss: 1.9037e-04
mean_absolute_error: 1.9037e-04
mean_squared_error: 3.6240e-08

...But horribly on balanced data with fewer features:
loss: 0.5981
mean_absolute_error: 0.5981
mean_squared_error: 0.8070
accuracy: 0.4788

The same problem: Only predicting no casualties

Failed Models: Neural Net

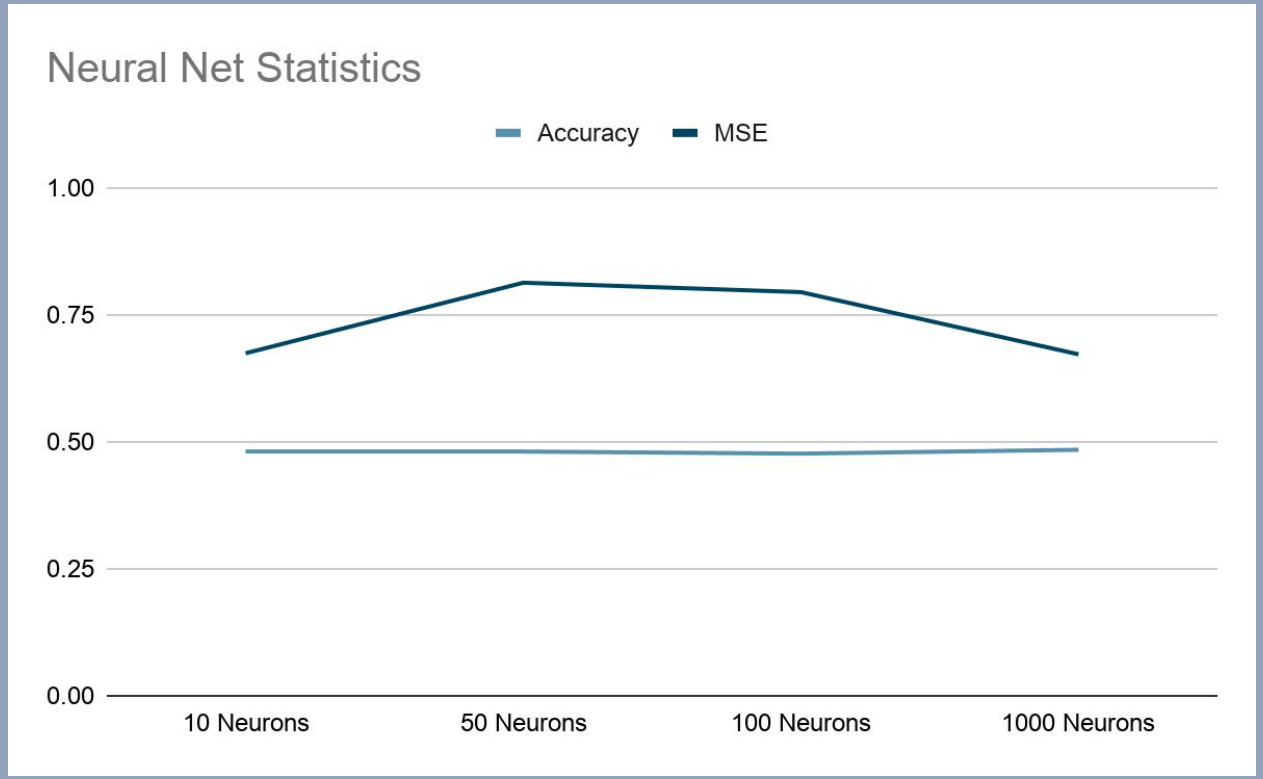

Couldn't improve accuracy or MSE

- ***Increasing neurons***
- ***Changing activation functions***
- ***Adding dropout layers***

Best accuracy ~ 0.5
Best MSE ~ 0.65

Final Model: KNN Classifier

Trained a KNN Classifier:

- *Did an 80/20 train_test_split*
- *Transformed CASUALTIES to a binary target*
- *Chose 6 best features from MI scores*

Mean Accuracy Score: 0.51

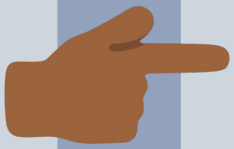

Attempted to up the score by:

- ***Hyperparam search (GridSearchCV, n_neighbors = 24)***
- ***Performed Cross Validation***
- ***PCA***
- ***NCA***
- ***Made derived columns (ex. "rush_hour" from "date" & "time")***
- ***Tried KNN, SGD***

... After all this, the best Mean Accuracy Score was: 0.65

Final Model: KNN Classifier

	Predicted false	Predicted true
Actual false	646	116
Actual true	408	321

	precision	recall	f1-score	support
0.0	0.61	0.85	0.71	762
1.0	0.73	0.44	0.55	729
accuracy			0.65	1491
macro avg	0.67	0.64	0.63	1491
weighted avg	0.67	0.65	0.63	1491

Product Demonstration

Lessons Learned

- **The data set may not have the best information for predicting how bad a crash will be. Research shows these things play major roles in fatal accidents:**
  - Speeding
  - Driver behavior (distracted driving, seat belts)
  - Car model
  - Head on collisions (typically caused by extreme driving errors/negligence)
- **We underestimated how important data cleaning, feature selection, and data transforming is to ML. Especially with messy real-life government datasets**
- **Our initial choice of data set was not ideal for our goal!**

Final Thoughts

- We're especially proud of how well we wrangled the dataset
- Although the classifier isn't performing amazingly well, it's getting over 0.5, which shows some predictive ability
- Our preprocessing techniques were successful since we did see improvement in our model after using them.
- We initially had really good results with regression and the NN, but we were able identify why the models weren't actually useful

Drive Safe!

Group 1: Unsupervised Twitter Robots

Dagny Stahl ◆ Luke Valentine ◆ Elise Renwick